

▼ SGD FOR LINEAR REGRESSION ON BOSTON HOME PRICE DATASET

▼ BOSTON HOME PRICE DATASET

Number of Instances 506

Number of Attributes 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

Attribute Information (in order)

CRIM - per capita crime rate by town

ZN - proportion of residential land zoned for lots over 25,000 sq.ft.

INDUS - proportion of non-retail business acres per town.

CHAS - Charles River dummy variable (1 if tract bounds river; 0 otherwise)

NOX - nitric oxides concentration (parts per 10 million)

RM - average number of rooms per dwelling

AGE - proportion of owner-occupied units built prior to 1940

DIS - weighted distances to five Boston employment centres

RAD - index of accessibility to radial highways

TAX - full-value property-tax rate per \$10,000

PTRATIO - pupil-teacher ratio by town

B - $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town

LSTAT - % lower status of the population

MEDV - Median value of owner-occupied homes in \$1000's (**prices of the house**)

```
1 from sklearn.datasets import load_boston
2
3 boston=load_boston()
```

```
1 boston.data.shape
```

```
↳ (506, 13)
```

```
1 boston.feature_names
```

```
↳ array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',
        'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7')
```

```
1 print(boston.DESCR)
```

```
↳
```

```
.. _boston_dataset:
```

Boston house prices dataset

****Data Set Characteristics:****

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(Bk - 0.63)^2$ where Bk is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity',
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth Internati

```
1 import numpy as np
2 from sklearn.preprocessing import StandardScaler
3 import pandas as pd
4 import warnings
5 warnings.filterwarnings('ignore')
```

```
1 boston_dataset=pd.DataFrame(data=boston.data)
```

```
1 boston_dataset.head(3)
```



	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03

```
1 boston_dataset.columns=boston.feature_names
```

```
1 boston_dataset.head(3)
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03

```
1 boston_dataset['MEDV']=boston.target
```

```
1 boston_dataset
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2
...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.67	22.4
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.08	20.6
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.64	23.9
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.48	22.0
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.88	11.9

506 rows × 14 columns

```
1 house_price=boston_dataset['MEDV']
2 boston_dataset = boston_dataset.drop("MEDV", axis = 1)
```

```
1 from sklearn.model_selection import train_test_split
2
3 X_Train,X_Test,Y_Train,Y_Test = train_test_split(boston_dataset,house_price,test_size=0.2,random_state=0)
```

▼ Standarize the data,

```
1 X_Train.shape
```

```
(404, 13)
```

```
1 from sklearn.preprocessing import StandardScaler
2
```

```

3 scalar = StandardScaler()
4 X_Train = scalar.fit_transform(np.array(X_Train))
5 X_Test = scalar.transform(np.array(X_Test))

```

```
1 manual_training = pd.DataFrame(data=X_Train)
```

```
1 manual_training
```

```

↳

```

	0	1	2	3	4	5	6	7	8	9	
0	-0.372574	-0.499608	-0.704925	3.664502	-0.424879	0.935678	0.693669	-0.437218	-0.162242	-0.561656	-0.48
1	-0.397099	-0.499608	-0.044878	-0.272888	-1.241859	-0.491181	-1.835528	0.730055	-0.624648	-0.573376	0.33
2	-0.402693	0.771168	-0.886760	-0.272888	-1.111490	0.857849	-0.524621	1.234769	-0.393445	-0.602677	-0.84
3	-0.405769	0.029882	-0.465819	-0.272888	-0.277127	-0.417676	-0.086464	0.861527	-0.509046	-0.538216	-1.48
4	2.774932	-0.499608	0.998884	-0.272888	1.070021	-1.438097	0.715042	-1.021528	1.687378	1.542121	0.79
...
399	-0.381700	-0.499608	-0.535975	-0.272888	-0.546557	-0.855823	0.187829	0.457582	-0.509046	-0.678858	0.51
400	-0.405972	1.406555	-1.100094	-0.272888	-1.033268	1.262847	-1.522051	1.303922	-0.509046	-0.028387	-1.48
401	-0.398056	-0.499608	-0.159419	-0.272888	-0.077228	-0.404705	0.483496	-0.492223	-0.393445	0.170857	-0.30
402	-0.388424	-0.499608	-0.603269	-0.272888	-0.937664	-0.391733	0.586801	0.923558	-0.740249	-0.995304	-0.25
403	-0.399513	-0.499608	-1.012756	-0.272888	-0.398805	-1.051836	0.693669	-0.565033	-0.509046	-0.626117	-0.84

404 rows × 13 columns

```
1 type(manual_training)
```

```
↳ pandas.core.frame.DataFrame
```

```
1 manual_training.columns=boston.feature_names
```


```
1 manual_training['MEDV']=Y_Train
```

```

1 from sklearn.linear_model import SGDRegressor
2 import seaborn as sns
3 from sklearn.metrics import mean_squared_error
4 import matplotlib.pyplot as plt
5
6 def sklearn_SGDRegressor(X_Train,Y_Train,X_Test,Y_Test,alpha,learning_rate,n_iterations):
7
8     linear_regression=SGDRegressor(alpha=alpha,learning_rate=learning_rate,n_iter_no_change=n_iterations)
9     linear_regression.fit(X_Train,Y_Train)
10
11     Y_Predicted=linear_regression.predict(X_Test)
12
13     plt.scatter(Y_Test,Y_Predicted)
14     plt.title('scatter plot between actual y and predicted y')
15     plt.xlabel('Actual Y')
16     plt.ylabel('Predicted Y')
17     plt.grid(b=True, linewidth=0.5)
18     plt.show()
19
20     return Y_Predicted,n_iterations,linear_regression.coef_, linear_regression.intercept_

```

```
1 y_predicted,n_iter,weight_sgd,intercept_sgd=sklearn_SGDRegressor(X_Train,Y_Train,X_Test,Y_Test,0.0001,'constant',1
```



```
1 sgd_error=mean_squared_error(Y_Test,y_predicted)
2 print('Mean Square Error(MSE)=', sgd_error)
3 print('Number of iteration=', n_iter)
```

Mean Square Error(MSE)= 32.80756905086225
Number of iteration= 1

```
1 print('W[j+1] :', weight_sgd )
2 print('b[j+1] :', intercept_sgd )
```

W[j+1] : [-0.81096515 1.0039064 0.02367056 0.54872808 -1.25903411 3.04495628
0.08831497 -2.72093591 1.4928305 -0.85047471 -2.50149652 1.23202212
-3.70111891]
b[j+1] : [22.56270378]

```
1 manual_training = manual_training.drop("MEDV", axis = 1)
```

▼ Using Manual Approach to find the Linear Regression

```
1 def find_differential_L_wrt_W(X_Data,Y_Data,W_tranpose,K,b,d):
2
3     # Calculate Differential of L w.r.t W
4     sum_diff = np.zeros(shape=(1,d))
5     for i in range(K):
6         second_term = (Y_Data.iloc[i] - b - np.dot(W_tranpose,X_Data[i]))
7         final_term = X_Data[i] * second_term
8         sum_diff += final_term
9     sum_diff *= (-2/K)
10    return sum_diff
```

```
1 def find_differential_L_wrt_b(X_Data,Y_Data,W_Transpose,K,b,d):
2     #Calculate differential of L w.r.t b
3     sum_diff = 0
4     for i in range(K):
5         final_term = Y_Data.iloc[i] - b - np.dot(W_Transpose,X_Data[i])
6         sum_diff += final_term
7     sum_diff *= (-2/K)
8     return sum_diff
```

```
1 def manual_predict(X_Test,W,b):
2     Y_Predicted=[]
3
4     for i in range(len(X_Test)) :
5         #predict = (np.dot(W,X_Test[i]) + b)
```

```

6     y=np.asscalar(np.dot(W,X_Test[i]) + b)
7     Y_Predicted.append(y)
8     return np.array(Y_Predicted)

```

```

1 def draw_scatterPlot(Y_Test,Y_Predicted) :
2     plt.scatter(Y_Test,Y_Predicted)
3     plt.title('scatter plot between actual y and predicted y')
4     plt.xlabel('Actual Y')
5     plt.ylabel('Predicted Y')
6     plt.grid(b=True, linewidth=0.5)
7     plt.show()

```

```

1 def find_and_transform(X,Y,K,n_iterations,threshold,power_t,lr_rate) :
2     n,d=X.shape
3     b=0
4     weight = np.random.randn(1,d)
5     r,iter_ = 1,1
6     n=1
7     while(n<=n_iterations):
8         w_prev = weight
9         b_prev = b
10        X_data=X.sample(K)
11        X_new=X_data[X_data.columns[0:]].values
12        Y_new = Y
13
14        Diff_W = find_differential_L_wrt_W(X_new,Y_new,w_prev,K,b_prev,d)
15        Diff_b = find_differential_L_wrt_b(X_new,Y_new,w_prev,K,b_prev,d)
16
17        #Calculate Wj+1 & bj+1
18
19        Wj_plus_1 = w_prev - (r*Diff_W)
20        bj_plus_1= b_prev - (r*Diff_b)
21
22        if(iter_ >= n_iterations):
23            break
24
25        #r = r/2
26        iter_ +=1
27        n=n+1
28
29
30    print('*****')
31
32    print('W[j+1] :' , Wj_plus_1)
33    print('b[j+1] :' , bj_plus_1)
34    #print('r :' , r)
35    print('*****')
36    return Wj_plus_1,bj_plus_1

```

```
1 W_star,b_star=find_and_transform(manual_training,Y_Train,50,1,0.1,0.25,0.01)
```

```


↳ *****
W[j+1] : [[-2.81993685  1.82824696  2.49608589  1.78137557 -0.80908617 -4.67501745
 -1.6640677   0.30820965  2.08038961  1.41089452  1.94006647  3.70436801
 -2.89953773]]
b[j+1] : [42.01570198]
r : 1
*****

```

```

1 Y_Predicted=manual_predict(X_Test, W_star,b_star)
2 draw_scatterPlot(Y_Test,Y_Predicted)
3 manual_error=mean_squared_error(Y_Test,Y_Predicted)
4 print('Mean Square Error(MSE) Using Manual Approach =', manual_error)


```



Mean Square Error(MSE) Using Manual Approach = 543.8083421548782

▼ Using SGD Regressor for No Of Iterations = 100

```
1 y_predicted,n_iter,weight_sgd,intercept_sgd=sklearn_SGDRegressor(X_Train,Y_Train,X_Test,Y_Test,0.0001,'constant',1
```



```
1 sgd_error=mean_squared_error(Y_Test,y_predicted)
2 print('Mean Square Error(MSE)=', sgd_error)
3 print('Number of iteration=', n_iter)
```

```
☐ Mean Square Error(MSE)= 36.589875144670955
   Number of iteration= 100
```

```
1 print('W[j+1] :', weight_sgd )
2 print('b[j+1] :', intercept_sgd )
```

```
☐ W[j+1] : [-0.76946944  1.25748731 -0.21317185  0.73145965 -1.95077354  2.68340876
 -0.20734028 -2.65472201  2.01102926 -2.01512418 -2.3411727  0.7714277
 -3.70722231]
   b[j+1] : [22.36872328]
```


▼ Using Manual Linear Regression for No Of Iterations = 100

```
1 W_star,b_star=find_and_transform(manual_training,Y_Train,50,100,0.1,0.25,0.01)
2 Y_Predicted=manual_predict(X_Test, W_star,b_star)
3 draw_scatterPlot(Y_Test,Y_Predicted)
4 manual_error=mean_squared_error(Y_Test,Y_Predicted)
5 print('Mean Square Error(MSE) Using Manual Approach =', manual_error)
```

```

*****
W[j+1] : [[-2.68493813  0.56919869  0.9558448  -2.27750919 -3.72175452 -1.67978081
-2.52142237  3.16574685  3.80873151  6.9708999  8.99502337 -4.45016381
-1.40969299]]
b[j+1] : [43.98986384]
r : 1
*****


```



Mean Square Error(MSE) Using Manual Approach = 953.3641591193748

▼ Using SGD Regressor for No Of Iterations = 1000

```
1 y_predicted,n_iter,weight_sgd,intercept_sgd=sklearn_SGDRegressor(X_Train,Y_Train,X_Test,Y_Test,0.0001,'constant',1
```



```

1 sgd_error=mean_squared_error(Y_Test,y_predicted)
2 print('Mean Square Error(MSE)=', sgd_error)
3 print('Number of iteration=', n_iter)

```

```

Mean Square Error(MSE)= 33.49280515235464
Number of iteration= 1000

```

```

1 print('W[j+1] :', weight_sgd )
2 print('b[j+1] :', intercept_sgd )

```

```

W[j+1] : [-0.95180992  1.05603303 -0.06516896  0.59334444 -2.03775517  3.24860141
 0.29953031 -2.88869107  2.02163147 -2.09274246 -2.64764731  0.49126429
-3.89843943]
b[j+1] : [22.3750692]

```

▼ Using Manual Regression for No Of Iterations = 1000


```


1 W_star,b_star=find_and_transform(manual_training,Y_Train,50,1000,0.1,0.25,0.01)
2 Y_Predicted=manual_predict(X_Test, W_star,b_star)
3 draw_scatterPlot(Y_Test,Y_Predicted)
4 manual_error=mean_squared_error(Y_Test,Y_Predicted)
5 print('Mean Square Error(MSE) Using Manual Approach =', manual_error)

```

```

*****
W[j+1] : [[-6.33871176 -3.66884595 -3.22007523  4.79155866  2.11014711  0.70846533
 -0.02835366  1.16817197 -0.24489098 -0.2492538  -3.90611166 -3.38394822
 -3.8233088  ]]
b[j+1] : [41.81694397]
*****

```




Mean Square Error(MSE) Using Manual Approach = 522.542501849131

▼ Using SGD Regressor for No Of Iterations = 10000

```

1 y_predicted,n_iter,weight_sgd,intercept_sgd=sklearn_SGDRegressor(X_Train,Y_Train,X_Test,Y_Test,0.0001,'constant',1

```



```

1 sgd_error=mean_squared_error(Y_Test,y_predicted)
2 print('Mean Square Error(MSE)=', sgd_error)
3 print('Number of iteration=', n_iter)

```

```

Mean Square Error(MSE)= 34.5350056094487
Number of iteration= 10000

```

```

1 print('W[j+1] :', weight_sgd )
2 print('b[j+1] :', intercept_sgd )

```

```



```

```
W[j+1] : [-1.00132737e+00  9.95600160e-01  8.65082611e-04  7.34974995e-01
-1.99042815e+00  1.95265621e+00 -5.56339997e-02 -2.61382905e+00
 2.13846077e+00 -1.78132803e+00 -2.12006519e+00  9.32669169e-01
-3.67285574e+00]
b[j+1] : [22.27394321]
```

▼ Using Manual Regression for No Of Iterations = 10000

```
1 W_star,b_star=find_and_transform(manual_training,Y_Train,50,10000,0.1,0.25,0.01)
2 Y_Predicted=manual_predict(X_Test, W_star,b_star)
3 draw_scatterPlot(Y_Test,Y_Predicted)
4 manual_error=mean_squared_error(Y_Test,Y_Predicted)
5 print('Mean Square Error(MSE) Using Manual Approach =', manual_error)
```

```
*****
W[j+1] : [[-3.23708485 -4.76309083 -2.38928386  0.86176101 -1.53611454 -0.70770322
 3.66204276  2.23203837 -4.93908421 -4.1641031 -6.15103513  2.59545455
-1.32903266]]
b[j+1] : [42.57705016]
*****
```



Mean Square Error(MSE) Using Manual Approach = 663.7141338124138

▼ COMPARISON TABLE

```
1 from prettytable import PrettyTable
2 x= PrettyTable()
3 x.field_names = ["Algorithm" , "Alpha", "No of Iterations","MSE", "B_Star"]
4 x.add_row(["SGD Regressor",0.0001,1,32.80,22.56])
5 x.add_row(["Manual SGD",0.0001,1,543.80,42.01])
6 x.add_row(["SGD Regressor",0.0001,100,36.58,22.36])
7 x.add_row(["Manual SGD",0.0001,100,953.80,44.98])
8 x.add_row(["SGD Regressor",0.0001,1000,33.49,22.37])
9 x.add_row(["Manual SGD",0.0001,1000,522.54,41.81])
10 x.add_row(["SGD Regressor",0.0001,10000,34.53,22.27])
11 x.add_row(["Manual SGD",0.0001,10000,663.71,42.57])
12 print(x)
```



Algorithm	Alpha	No of Iterations	MSE	B_Star
SGD Regressor	0.0001	1	32.8	22.56
Manual SGD	0.0001	1	543.8	42.01
SGD Regressor	0.0001	100	36.58	22.36
Manual SGD	0.0001	100	953.8	44.98
SGD Regressor	0.0001	1000	33.49	22.37
Manual SGD	0.0001	1000	522.54	41.81
SGD Regressor	0.0001	10000	34.53	22.27
Manual SGD	0.0001	10000	663.71	42.57

Looks like the Mean Square Error(MSE) getting decrease when No of Iterations are increasing. Also it is 50

Note : I have not reduced the r to $r/2$ for each iterations. I tried this approach but then the MSE and B_Star getting increased drastically.